

```
In [1]: import pandas as pd
import numpy as np

In [2]: telco = pd.read_csv("WA_Fn-UseC_-Telco-Customer-Churn.csv")
type(telco)

Out[2]:pandas.core.frame.DataFrame

In [3]: telco.head()

Out[3]:
```

	customerID	gender	SeniorCitizen	Partner	Dependents	tenure	PhoneService	MultipleLines	InternetService	OnlineSecurity	...	DeviceProtection
0	7590-VHVEG	Female	0	Yes	No	1	No	No phone service	DSL	No	...	No
1	5575-GNVDE	Male	0	No	No	34	Yes	No	DSL	Yes	...	Yes
2	3668-QPYBK	Male	0	No	No	2	Yes	No	DSL	Yes	...	No
3	7795-CFOCW	Male	0	No	No	45	No	No phone service	DSL	Yes	...	Yes
4	9237-HQITU	Female	0	No	No	2	Yes	No	Fiber optic	No	...	No

5 rows x 21 columns



```
In [ ]: 1. The target variable in this dataset is "Churn." It is the variable that you are trying to predict. Specifically, you want to determine whether a customer v

In [ ]: 2. This is a binary classification problem. You are trying to classify customers into two categories: those who will churn and those who will stay.

In [4]: column_names = telco.columns
data_types = telco.dtypes

print("Column Names:")
print(column_names)

print("\nData Types:")
print(data_types)

Column Names:
Index(['customerID', 'gender', 'SeniorCitizen', 'Partner', 'Dependents',
      'tenure', 'PhoneService', 'MultipleLines', 'InternetService',
      'OnlineSecurity', 'OnlineBackup', 'DeviceProtection', 'TechSupport',
      'StreamingTV', 'StreamingMovies', 'Contract', 'PaperlessBilling',
      'PaymentMethod', 'MonthlyCharges', 'TotalCharges', 'Churn'],
      dtype='object')

Data Types:
customerID      object
gender          object
SeniorCitizen   int64
Partner         object
Dependents      object
tenure          int64
PhoneService    object
MultipleLines   object
InternetService object
OnlineSecurity  object
OnlineBackup    object
DeviceProtection object
TechSupport     object
StreamingTV     object
StreamingMovies object
Contract        object
PaperlessBilling object
PaymentMethod   object
MonthlyCharges  float64
TotalCharges    object
Churn           object
dtype: object

In [5]: missing_values = telco.isnull().sum()
print(missing_values)
```

```

customerID      0
gender          0
SeniorCitizen  0
Partner        0
Dependents     0
tenure         0
PhoneService   0
MultipleLines  0
InternetService 0
OnlineSecurity 0
OnlineBackup   0
DeviceProtection 0
TechSupport    0
StreamingTV    0
StreamingMovies 0
Contract       0
PaperlessBilling 0
PaymentMethod  0
MonthlyCharges 0
TotalCharges   0
Churn          0
dtype: int64
In [6]: telco = telco.dropna()

In [7]: duplicates = telco[telco.duplicated()]
        print(duplicates)

Empty DataFrame
Columns: [customerID, gender, SeniorCitizen, Partner, Dependents, tenure, PhoneService, MultipleLines, InternetService, OnlineSecurity, OnlineBackup, DeviceProtection, TechSupport, StreamingTV, StreamingMovies, Contract, PaperlessBilling, PaymentMethod, MonthlyCharges, TotalCharges, Churn]
Index: []

[0 rows x 21 columns]
In [8]: telco = telco.drop_duplicates()

In [ ]: Use methods like label encoding or one-hot encoding if there are category qualities that need to be converted into numerical values. The categorical fe

In [11]: churn_counts = telco['Churn'].value_counts()
         print(churn_counts)

Churn
No    5174
Yes   1869
Name: count, dtype: int64

In [ ]: An outlier is an observation that lies an abnormal distance from other values in a dataset. For example, a very high or extremely low income compared

In [ ]: You can use techniques like Z-score or IQR (Interquartile Range) to detect outliers and handle them. Depending on the specific analysis and how outliers

In [17]: import seaborn as sns

In [18]: correlation = telco['Churn'].replace({'No':0, 'Yes':1}).corr(telco['gender'].replace({'Male':0, 'Female':1}))
         print(correlation)


0.008612095078997867
There is no pattern

In [19]: churn_percent = (telco['Churn'].value_counts() / len(telco)) * 100
         print(churn_percent)

Churn
No    73.463013
Yes   26.536987
Name: count, dtype: float64
In [20]: import matplotlib.pyplot as plt
         plt.figure(figsize=(8, 6))
         plt.hist (telco['tenure'], bins=20, color='red')
         plt.xlabel("Tenure (Months)")
         plt.ylabel("Number of Customers")
         plt.title('Distribution of customers Tenure')
         plt.show()

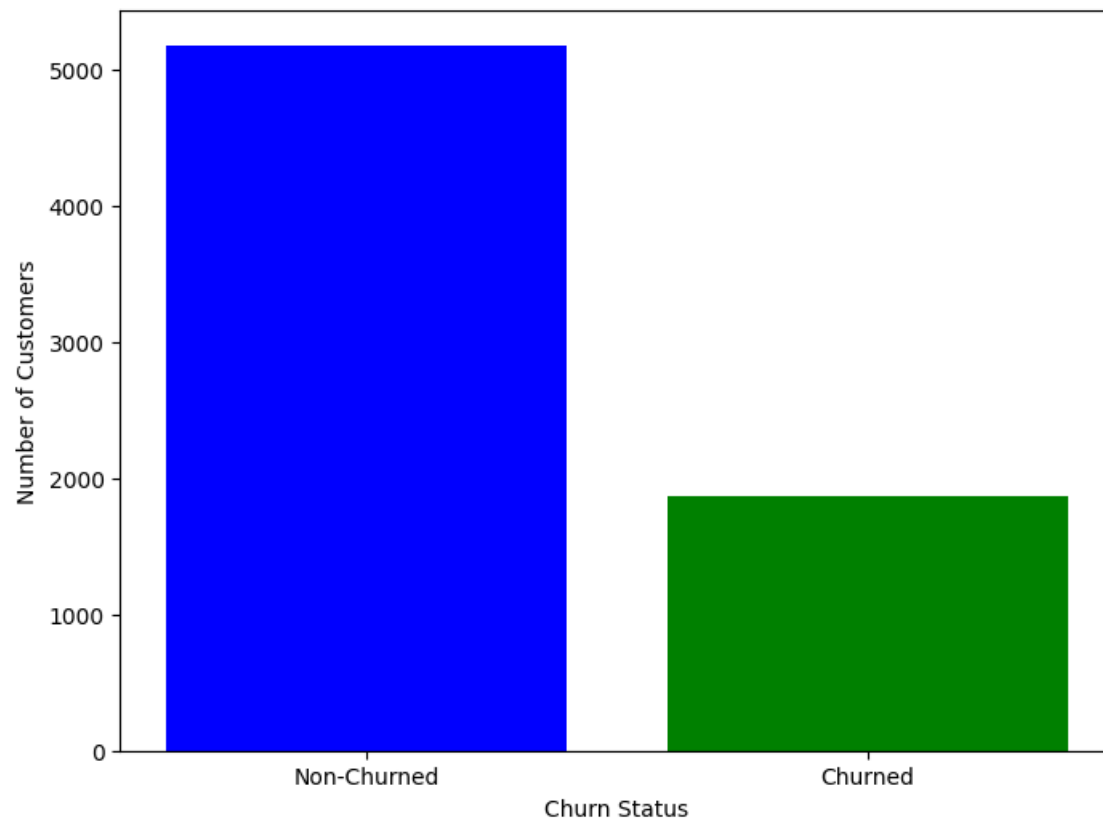
```

Distribution of customers Tenure



```
In [21]: plt.figure(figsize=(8, 6))
plt.bar(churn_counts.index, churn_counts.values, color=['blue', 'green'])
plt.xlabel('Churn Status')
plt.ylabel('Number of Customers')
plt.title('Churned vs. Non-Churned Customers')
plt.xticks([0,1], ['Non-Churned', 'Churned'])
plt.show()
```


Churned vs. Non-Churned Customers



```
In [22]: churn_percent = (telco['Churn'].value_counts() / len(telco)) * 100
plt.figure(figsize=(8, 6))
plt.pie(churn_percent, labels=['Non-Churned', 'Churned'], autopct='%1.1f%%', colors=['yellow', 'lightgreen'])
plt.title('Percentage of Churned vs. Active Customers')
plt.show()
```

Percentage of Churned vs. Active Customers

Non-Churned



26.5%

Churned